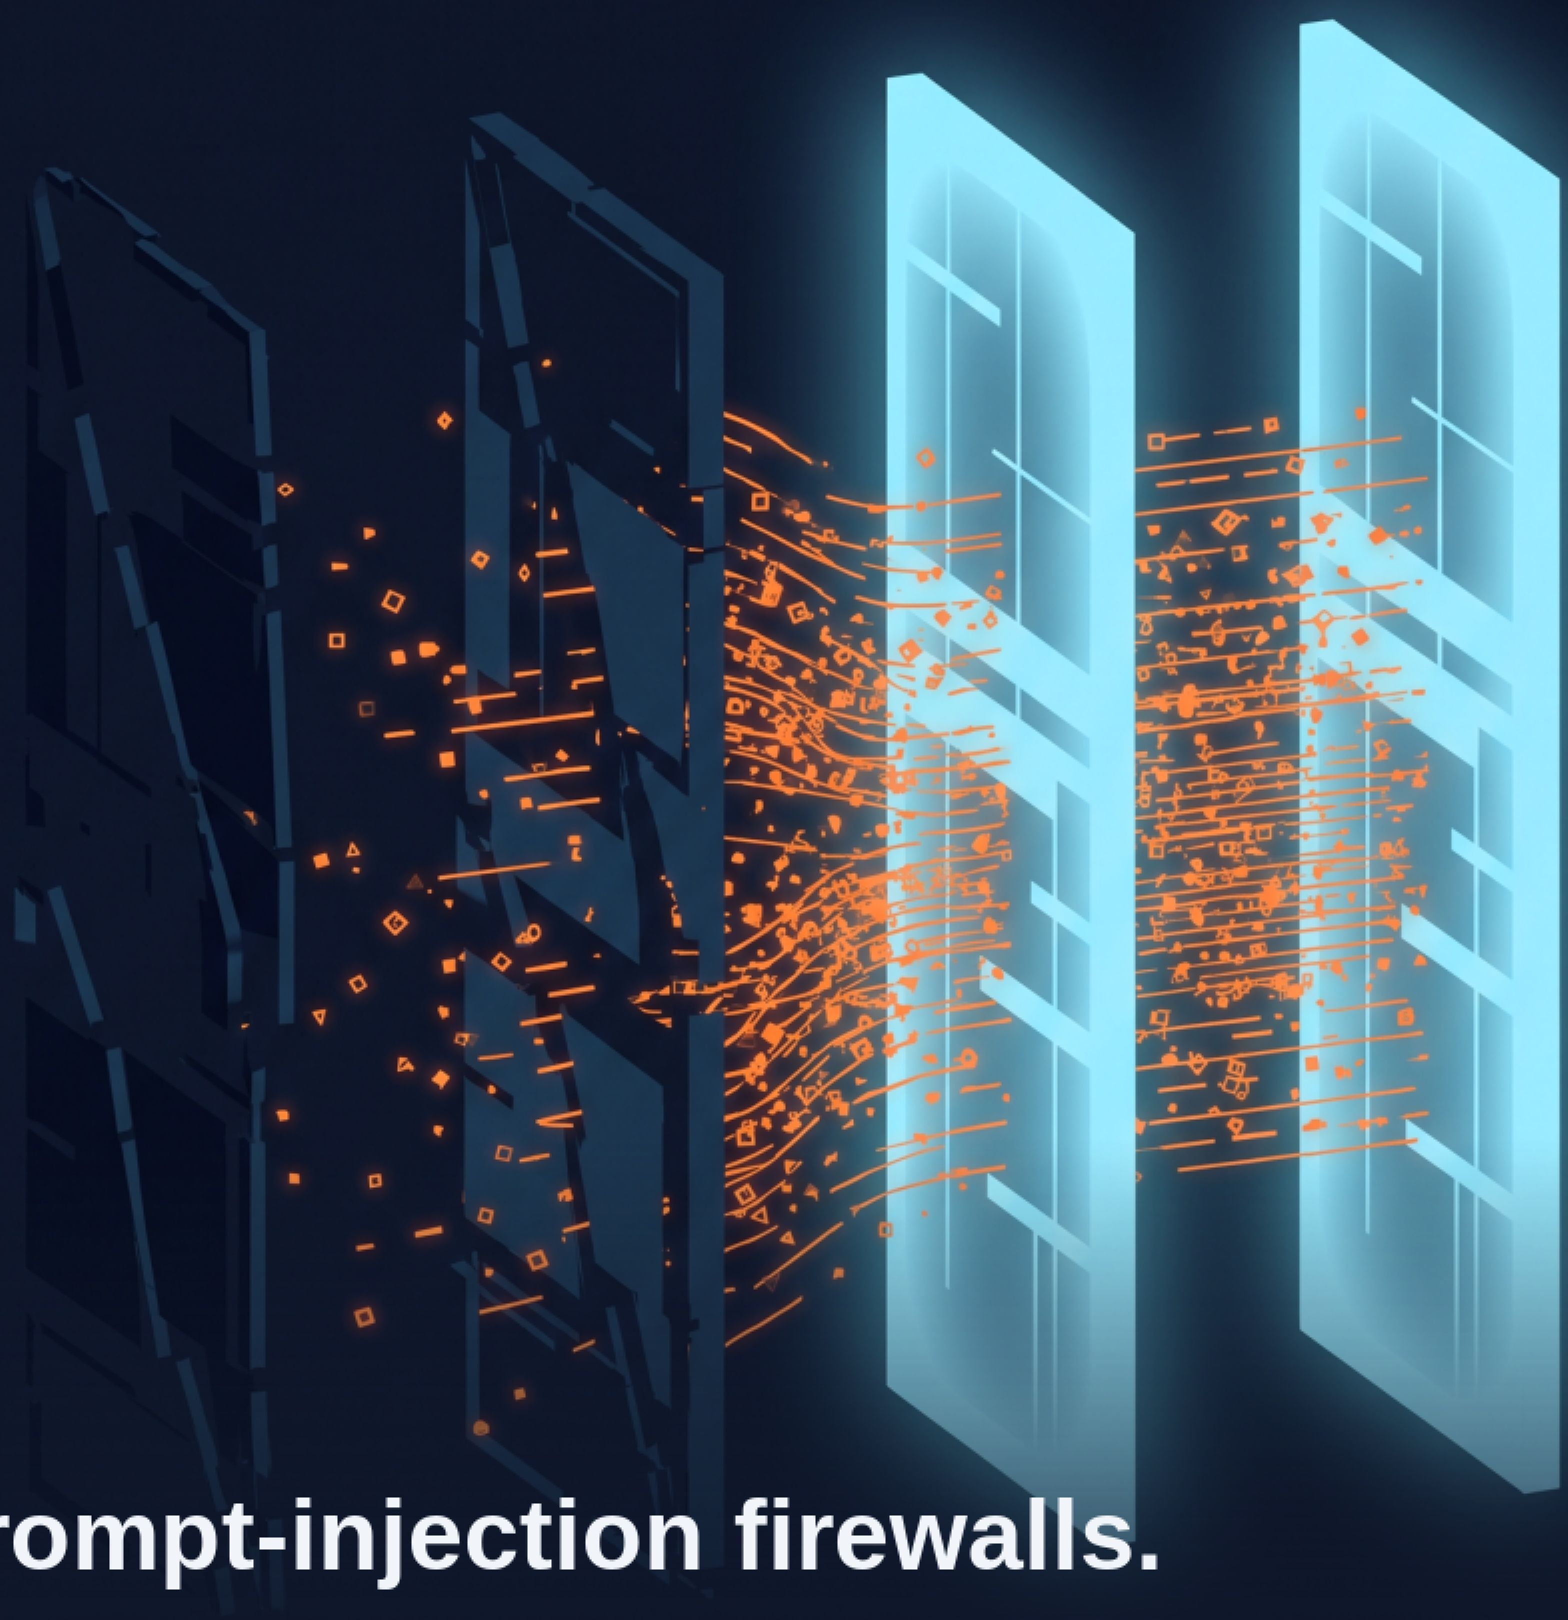

4 prompt-injection firewalls.
Only 2 would even run.

Both caught every attack. Only one is deployable.

Two never started

Prompt Guard 2

Meta / Llama-gated

The model repo is gated behind Meta's license. Access request still pending — couldn't even download the weights.

Vigil

Won't build offline

Died after 3 rebuilds and 6 separate blockers: missing YARA, dep conflicts, a VectorDB built unconditionally at startup.

Both caught all 6

I sent an escalating battery through the two guards that ran:

Plain direct injection
caught by both

Spaced / homoglyph obfuscation
caught by both

Base64-encoded payload
caught by both

Indirect — hidden in a retrieved doc
caught by both

On detection alone: a perfect tie.
Which is exactly where most comparisons stop.

The tie-breaker

Then the test that actually matters — the benign set, and the clock:

LLM Guard

33%
over-blocked (false positives)

449 ms
per call (p50 latency)

Rebuff

67%
over-blocked (false positives)

3.7 s
per call (p50 latency)

Rebuff blocks two-thirds of harmless traffic at 3.7s a call. That's not a firewall — it's a **denial-of-service on your own users.**

The verdict

1	LLM Guard	8.1
2	Rebuff	6.7

Prompt Guard 2 and Vigil — not evaluated (never ran).

A single classifier is a speed bump, not a wall.

Detection scores all look great — until you measure what they block by mistake and how long they take. Layer your defenses; the **firewall is never your only line.**

Full teardown + reproducible lab in the comments.