


Four prompt-injection firewalls, tested.


Only **2**
would even run.

THE FIELD – FOUR CONTENDERS



LLM Guard

RAN



Rebuff

RAN



Prompt Guard 2

NEVER RAN



Vigil

NEVER RAN

Both that ran caught every attack. **Only one is deployable.**

Two never made the start line.

01 / 02

DID NOT RUN

Prompt Guard 2

META • LLAMA-GATED

The model repo is gated behind Meta's license. My access request is *still pending* — I couldn't even download the weights.

02 / 02

DID NOT RUN

Vigil

WON'T BUILD OFFLINE

Died after 3 rebuilds and 6 separate blockers: missing YARA, dependency conflicts, a VectorDB built unconditionally at startup.

FINDING #1 Half the best-known field doesn't install cleanly — the one thing no vendor puts on a slide.

Both survivors caught all six.

An escalating battery, run through the two guards that started.


	LLM GUARD	REBUFF
Plain direct injection	✓	✓
Spaced / homoglyph	✓	✓
Base64-encoded payload	✓	✓
Indirect — hidden in a doc	✓	✓

6/6 - 6/6

A perfect tie.
On detection alone – which is exactly where most comparisons stop.

Then the test that matters.

The benign set – harmless prompts, and the clock.



Blocking two-thirds of harmless traffic at 3.7s a call isn't a firewall. It's a **denial-of-service on your own users.**

The verdict.

01	LLM Guard DEPLOYABLE • 449ms • 33% FP	8.1
02	Rebuff UNUSABLE • 3.7s • 67% FP	6.7

Prompt Guard 2 & Vigil – not evaluated (never ran).

A single classifier is a speed bump, not a wall.

Detection scores all look great until you measure what they block by mistake and how long they take. So you layer: a fast, low-false-positive gate at the door, output checks behind the model, canary tokens, least-privilege. **The firewall is never your only line.**

Full teardown + reproducible lab in the comments →